

# COVID-19 Retweet Prediction Report

August 5, 2021

## 1 Team Members

1. Aaron Khoo
2. Calvin Yusnoveri (1002911)
3. Amarjyot Kaur Narula
4. Joseph Chng (1003811)

## 2 Task Description

As COVID-19 impacts our daily routine and changed the norms that we accepted prior to the pandemic, there are some interests in quantifying its impacts on the global stage. One way to measure such impact is to monitor the explosion of activity in social media usage such as Twitter and Youtube, where most people who are not able to move freely, shares their thoughts through such platforms. Twitter in particular, provides a platform that allows the users to post their thoughts in a succinct manner and add hashtags or mentions to increase the tweet's exposure on the platform. Our task will thus be to predict the number of retweets a tweet that is COVID-19 related will have using the TweetsCOV-19 dataset.

At the end of this project, we created a Linear Regression Model and an interactive GUI to predict retweet count based on input parameters found in TweetsCOV-19 dataset. Custome values can also be input into the model.

### 3 Dataset Description

The dataset that is used for this project is obtained from the COVID-19 Retweet Prediction Challenge. For this prediction model, we used Part 2 dataset that can be obtained from this website <https://data.gesis.org/tweetscov19/#dataset>. This dataset consists of tweets that is COVID-19 related from the month of May 2020.

From the dataset, there are different features for the tweet data that we obtain and the feature description are as follows:

1. Tweet Id: Long. Unique ID for a specific tweet
2. Username: String. Username of the user that published the tweet which is encrypted for privacy.
3. Timestamp: Format ( "EEE MMM dd HH:mm:ss Z yyyy" ). Specific time and date of the tweet
4. #Followers: Integer. Number of followers of the Twitter user who posted the tweet.
5. #Friends: Integer. Number of friends that the Twitter user who posted the tweet.
6. #Retweets: Integer. Number of retweets that the tweet has obtained and is the label for this project.
7. #Favorites: Integer. Number of favorites for the tweet
8. Entities: String. The entities of the tweet is obtained by aggregating the original text. Every annotated entity will then have its produced score from FEL library. Each entity is separated by char ":" to store the entity in this form "original\_text:annotated\_entity:score;". Each entity is separated from another entity by char ";". Any tweet that has no corresponding entities will be stored as "null;".
9. Sentiment: String. SentiStrength produces a score for positive (1 to 5) and negative (-1 to -5) sentiment. The two sentiments are splitted by whitespace char " ". Positive sentiment was stored first and followed by negative sentiment (i.e. "2 -1").
10. Mentions: String. Contains mentions and concatenate them with whitespace char " ". If there is no mention, it is stored as "null;".
11. Hashtags: String. Contains hashtags and concatenate the hashtags with whitespace char " ". If there is no hashtag, it is stored as "null;".
12. URLs: String. Contains URLs and concatenate the URLs using ":-:". If there is no URL, it is stored as "null;".

```
[1]: # imports

import logging
import requests
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from gensim.models import Word2Vec
```

```
c:\users\calvin yusnoveri\appdata\local\programs\python\python36\lib\site-
packages\gensim\similarities\__init__.py:15: UserWarning: The
gensim.similarities.levenshtein submodule is disabled, because the optional
Levenshtein package <https://pypi.org/project/python-Levenshtein/> is
```

unavailable. Install Levenshtein (e.g. `pip install python-Levenshtein`) to suppress this warning.  
 warnings.warn(msg)

```
[5]: header = [
    "Tweet Id",
    "Username",
    "Timestamp",
    "#Followers",
    "#Friends",
    "#Retweets",
    "#Favorites",
    "Entities",
    "Sentiment",
    "Mentions",
    "Hashtags",
    "URLs"]

data = pd.read_csv("./data/TweetsCOVID19_052020.tsv.gz", compression='gzip',
    ↪names=header, sep='\t', quotechar='')
data.head(5)
```

```
[5]:
```

	Tweet Id	Username \
0	1255980348229529601	fa5fd446e778da0acba3504aeab23da5
1	1255981220640546816	547501e9cc84b8148ae1b8bde04157a4
2	1255981244560683008	840ac60dab55f6b212dc02dcbe5dfbd6
3	1255981472285986816	37c68a001198b5efd4a21e2b68a0c9bc
4	1255981581354905600	8c3620bdfb9d2a1acfd2412c9b34e06

  

	Timestamp	#Followers	#Friends	#Retweets \
0	Thu Apr 30 22:00:24 +0000 2020	29697	24040	0
1	Thu Apr 30 22:03:52 +0000 2020	799	1278	4
2	Thu Apr 30 22:03:58 +0000 2020	586	378	1
3	Thu Apr 30 22:04:52 +0000 2020	237	168	0
4	Thu Apr 30 22:05:18 +0000 2020	423	427	0

  

	#Favorites	Entities	Sentiment \
0	0	null;	1 -1
1	6	null;	1 -1
2	2	null;	2 -1
3	0	null;	1 -1
4	0	i hate u:I_Hate_U:-1.8786140035817729;quaranti...	1 -4

  

	Mentions	Hashtags \
0	null;	Opinion Next2blowafrika thoughts
1	null;	null;
2	null;	null;

```

3     null;                                null;
4     null;                                null;

                                URLs
0                                null;
1                                null;
2     https://www.bbc.com/news/uk-england-beds-bucks...
3     https://lockdownsceptics.org/2020/04/30/latest...
4                                null;

```

## 4 Preprocessing

In order to train our prediction model, we have also done some preprocessing of the features that are available in the dataset. All these changes to the raw features allow us to link these processed features to the final retweet prediction in a more precise manner.

Clean Data (Final structure/form of data before it is fed into the model): 1. #Followers: Float. Log transformed:  $\log_{10}(x + 1)$  2. #Friends: Float. Log transformed:  $\log_{10}(x + 1)$  3. #Favorites: Float. Log transformed:  $\log_{10}(x + 1)$  4. Positive (Sentiment): Float. Scaled. 5. Negative (Sentiment): Float. Scaled. 6. Sentiment Disparity: Float. Scaled. 7. No. of Entities: Float. Log transformed:  $\log_{10}(x + 1)$  8. Day of Week: Float. One-Hot Vector. (7, ) Vector. 9. Time Int: Float. Log transformed:  $\log_{10}(x + 1)$  10. Hashtags Embedding: (25, ) Vector. 11. Mentions Embedding: (25, ) Vector.

In total, the input dimension (first layer) is  $8 + 7 + 25 + 25 = (65, )$ .

The target is: #Retweets: Integer. Log transformed:  $\log_{10}(x + 1)$

### 4.1 Hashtags & Mentions

Both Hashtags and Mentions are in the form of list of Strings separated by whitespace. Thus, in order to create tractable input for the model, embeddings are created for both the Hashtags and Mentions of size (25, ).

```

[19]: hashtag_embeddings = Word2Vec.load('./data/hashtag_embeddings')
      mention_embeddings = Word2Vec.load('./data/mention_embeddings')

```

```

hashtags_vocab = hashtag_embeddings.wv.index_to_key
mentions_vocab = mention_embeddings.wv.index_to_key

```

```

print(hashtags_vocab[:5]) # example of hashtags key
print(mentions_vocab[:5]) # example of mentions key

```

```

['COVID19', 'coronavirus', 'Covid_19', 'covid19', 'May']
['realDonaldTrump', 'PMOIndia', 'narendramodi', 'jaketapper', 'YouTube']

```

```

[21]: hashtags_example = 'COVID19'
      mentions_example = 'realDonaldTrump'

```

```
print(f"{hashtags_example} -> {hashtags_embeddings.wv[hashtags_example]}")
print(f"{mentions_example} -> {mentions_embeddings.wv[mentions_example]}")
```

```
COVID19 -> [ 0.10622272  0.26996937 -0.46450084  0.10561462 -0.5595082
0.26207525
 0.28835535  0.80339587  0.30626374 -0.13036335  0.8120623 -0.46314418
 0.20126966 -0.9723947 -1.0051426 -0.04809839  0.4593365  0.09532893
-0.21894015 -0.23557915  0.42107382  0.4622469  0.53460604 -0.589559
 0.6296402 ]
realDonaldTrump -> [ 0.5991303 -0.10410535  0.23690729 -0.23115875 -0.961905
-0.11418784
 0.12405131  0.4196795 -1.493182 -0.20270342  1.2276924 -1.3593616
-0.19556278  0.27365074  0.32451993  1.9415929 -0.20647514 -0.17526582
-0.69910485 -1.6436449  1.3161302 -0.17269903 -0.5424232  1.0386076
 1.062889 ]
```

These embeddings are trained over 5 epoch and only considers String symbols that occur at least 200 times to be relevant. This is done by passing the argument `min_count=200`, when training. The hashtags and mentions vocabularies are saved in: `data`.

Using these embeddings, both Hashtags and Mentions column are iterated over and converted into vectors of size (25, ). With these rules: - For those Hashtags/Mentions cells that contain `null`, 0 vector of size 25 is outputted - For those Hashtags/Mentions cells that contain String symbols that occur less than 200 times (hence, not in vocab), they're treated as `null` - For those Hashtags/Mentions cells that contain multiple String symbols, their embedding vectors are summed

#### 4.1.1 Effectiveness of Mentions & Hashtags and their Embeddings

Before attempting to create these embeddings, a quick exploration was done to check the relevance of these Mentions and Hashtags in predicting Retweet score.

An initial assumption is that, certain Hashtags or Mentions would correlate in higher Retweet score. But, a quick look of data seems to suggest little correlations as most high Retweet score have `null` Mentions and Hashtags.

```
[19]: sorted_by_retweets = data.sort_values(by='#Retweets', ascending=False)
sorted_by_retweets.head(10) # observe that most Mentions and Hashtags are null
↳ for top #Retweets
```

```
[19]:
```

	Tweet Id	Username \
1637862	1265465820995411973	0d4d9b3135ab4271ea36f4ebf8e9eae9
1208647	1266553959973445639	c9378a990def5939fb179e034a0d402e
1328169	1258750892448387074	1921c65230cd080c689dc82ea62e6e74
1736035	1263579286201446400	7c4529bc4da01f288b95cd3876b4da47
751238	1266546753182056453	32634ab407c86a56dde59551b3871c42
702118	1259975524581064704	69745f3009b864ba75b7d066ade0adba
1037044	1266738565641371648	71b9c38db144b44e4cbbda75c9fbf272
482286	1267066200049229824	56eb2d106e7611ab8bb76de07af8f318
1812643	1256657625334284292	6b7cc62c18b45d1eee1c34eb375e72a4

1401494 1260237550091935746 6b49e6ca36daebd1048d59b1459026ae

				Timestamp	#Followers	#Friends	#Retweets	\
1637862	Wed	May	27	02:12:17 +0000 2020	3317	3524	257467	
1208647	Sat	May	30	02:16:10 +0000 2020	18661	0	135818	
1328169	Fri	May	08	13:29:33 +0000 2020	83320	1753	88667	
1736035	Thu	May	21	21:15:52 +0000 2020	451	359	82495	
751238	Sat	May	30	01:47:31 +0000 2020	1545	874	66604	
702118	Mon	May	11	22:35:48 +0000 2020	6106969	726	63054	
1037044	Sat	May	30	14:29:43 +0000 2020	45941	4550	61422	
482286	Sun	May	31	12:11:37 +0000 2020	678	524	61038	
1812643	Sat	May	02	18:51:40 +0000 2020	778	694	60719	
1401494	Tue	May	12	15:57:00 +0000 2020	3704	1144	60650	

	#Favorites	Entities	\
1637862	845579	tear gas:Tear_gas:-1.688018296396458;	
1208647	363852	null;	
1328169	224288	mike pence:Mike_Pence:-0.6712149436851893;ppe:...	
1736035	225014	null;	
751238	193599	douche:Douche:-2.0041883604919835;	
702118	248214	null;	
1037044	100570	null;	
482286	101117	quarantine:Quarantine:-2.3096035868012508;	
1812643	213614	null;	
1401494	214508	flatten the curve:Flatten_the_curve:-1.6515462...	

	Sentiment	Mentions	Hashtags	URLs
1637862	1 -1	null;	null;	null;
1208647	1 -3	null;	null;	null;
1328169	1 -1	null;	null;	null;
1736035	1 -1	null;	null;	null;
751238	3 -1	null;	null;	null;
702118	1 -1	null;	null;	null;
1037044	1 -1	null;	null;	null;
482286	2 -1	null;	null;	null;
1812643	1 -1	null;	null;	null;
1401494	1 -1	null;	null;	null;

However, it is believed that there should at least be some value in including these Mentions and Hashtags even though such correlations are weak and not easily discernable. Thus, the embeddings are created regardless of the known weak correlation.

As for the embeddings themselves, based on similarity scores, they seem to be working well. For instance, the embedding are able to recognize `coronavirus` to be similar to `pandemic`, `COVID` and `virus` fairly confidently.

```
[20]: hashtag_embeddings.wv.most_similar(['coronavirus'])
```

```
[20]: [('virus', 0.7936981320381165),
      ('pandemic', 0.6945043802261353),
      ('COVID', 0.6876555681228638),
      ('corona', 0.6836724281311035),
      ('mask', 0.6728377342224121),
      ('trump', 0.6640805006027222),
      ('masks', 0.6489465832710266),
      ('covid', 0.6476311087608337),
      ('ClimateChange', 0.6469046473503113),
      ('COVID__19', 0.6411719918251038)]
```

## 4.2 Timestamp

## 4.3 Sentiment

## 4.4 Entities

The entities that are encapsulated in this dataset are aggregated from the original tweet text. This text will then go through a Fast Entity Linker query and find annotated text that can be found and set them as an entity for the text that we pass through. For every entity, it also has its corresponding log-likelihood confidence score which is used as a global threshold for linking.

For this project, we did some preprocessing of the raw entity data from the dataset. With the format of the entities for each tweet data being “original text:annotated text:score;original text:annotated text:score”, we will be able to get the number of entities that is found on each tweet. We thus split the entities column data and obtain the length of each entities list to get the number of entities for each tweet.

The column that contains the number of entities for each tweet will then undergo logarithmic transformation of the form  $\log(x+1)$  before they’re passed into the model. These log-transformed entity count data will then be used for the training of the prediction model and helps in creating a stronger linkage between the data and the number of retweets.

```
[12]: entities = data['Entities'].str.split(";")
      entity_no = []
      for ent in entities:
          ent.pop()
          if ent[0]=='null':
              entity_no.append(0)
          else:
              entity_no.append(len(ent))
      data['No. of Entities'] = entity_no
      # print(len(entities))
      data[['Entities', 'No. of Entities']].head(10)
```

```
[12]:
```

	Entities	No. of Entities
0	null;	0
1	null;	0
2	null;	0

3		0
4	i hate u:I_Hate_U:-1.8786140035817729;quaranti...	2
5	god forbid:God_Forbid:-1.2640735877261988;covi...	3
6	beijing:Beijing:-1.4222174822860647;covid 19:C...	3
7		0
8	stealth:Stealth_game:-2.646174787470186;	1
9	quarantine:Quarantine:-2.3096035868012508;	1

## 5 Model Architecture

We use ensemble methods (insert image): 1. 0-classifier (print out layer) 2. regression model (print out layer)

## 6 Results

Loss curve image. (Maybe save to .txt when training so can read independently and display with matplotlib independently.)

Accuracy on train and test set. (just run test.py on test set)

Validation images. (just screenshot gui)

## 7 Discussion

comparing with state of the art.

possible issues. possible improvements

## 8 GUI

### 8.1 Running the GUI

0. Download the dataset called `TweetsCOV19_052020.tsv.gz` and save it in `data` directory.  
Download from: <https://zenodo.org/record/4593502#.YQunN4gzZPY>
1. Open command line or terminal and navigate to the project folder.
2. Run `pip install -r requirements.txt`. Ensure that `PyQt5`, `gensim` and `torch` are installed among other things.
3. Navigate to `gui` with `cd gui`
4. Then, run `app.py` with `python app.py` or with your IDLE. Note: `app.py` must be executed from `./gui` NOT root!



**Retweet Prediction App**

Randomize	Predict	Predicted Retweet:
#Followers (int):	<input type="text"/>	-
#Friends (int):	<input type="text"/>	
#Favorites (int):	<input type="text"/>	True Retweet (if any):
Sentiment (str):	<input type="text"/>	
Datetime (ISO):	<input type="text"/>	-
Mentions:	<input type="text"/>	
Hashtags:	<input type="text"/>	No real datapoint referenced.
No. of Entities:	<input type="text"/>	

Opening the GUI may take a while (~3 mins) because it will load the dataset from `data/TweetsCOV19_052020.tsv.gz`. Thus, ensure that this file exist in `data` folder!

## 8.2 Using the GUI

1. Click **Randomize** Button to randomly load a data from dataset.
2. Click **Predict** Button to predict the rewteet.

**Retweet Prediction App**

Randomize	Predict	Predicted Retweet:
#Followers (int):	24890	17
#Friends (int):	730	
#Favorites (int):	43	True Retweet (if any):
Sentiment (str):	1 -1	
Datetime (ISO):	Mon May 04 15:13:03 +0000 2020	18
Mentions:	null;	
Hashtags:	MVSales	Data referenced. Index: 655122. Tweet Id: 1257327386581925889.
No. of Entities:	2	

Optionally: 3. Edit the line edits to create your own custom datapoint before predicting.

Note: - Fields like **Mentions** and **Hashtags** have their own vocab (consult: `data/hashtags_vocab.txt` and `data/mentions_vocab.txt` for valid values. Random values will just get ignored. - Each fields have their own format, like the ISO date string. Please

follow the format strictly, otherwise it will fail to be parsed. - Some fields can take in multiple values such as **Mentions** and **Hashtags**, in this case, use white space " " to delimit values.

Retweet Prediction App

Retweet Prediction App

Randomize

Predict

#Followers (int):

50000

#Friends (int):

200

#Favorites (int):

1000

Sentiment (str):

5 -1

Datetime (ISO):

Fri May 22 02:50:16 +0000 2020

Mentions:

WHO YouTube

Hashtags:

COVID19 WFH

No. of Entities:

2

Predicted Retweet:

207

True Retweet (if any):

-

No real data referenced. You are entering custom data.

## 9 Sources

1. Source code: <https://github.com/arglux/50021-ai-project>
2. Report:
3. Reference papers:

[ ]: